

MACS33002: Homework #3

Yujiao Song

Due Monday, Feb 17 by 5pm

```
library(gridExtra)
library(tidyverse)

## -- Attaching packages ----- tidyverse 1.3.0 --

## v ggplot2 3.2.1      v purrr   0.3.3
## v tibble  2.1.3      v dplyr  0.8.3
## v tidyr   1.0.2      v stringr 1.4.0
## v readr   1.3.1      v forcats 0.4.0

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::combine() masks gridExtra::combine()
## x dplyr::filter()  masks stats::filter()
## x dplyr::lag()     masks stats::lag()

library(mosaic)

## Loading required package: lattice

## Loading required package: ggformula

## Loading required package: ggstance

##
## Attaching package: 'ggstance'

## The following objects are masked from 'package:ggplot2':
##
##     geom_errorbarh, GeomErrorbarh

##
## New to ggformula? Try the tutorials:
##   learnr::run_tutorial("introduction", package = "ggformula")
##   learnr::run_tutorial("refining", package = "ggformula")
```

```

## Loading required package: mosaicData

## Loading required package: Matrix

##
## Attaching package: 'Matrix'

## The following objects are masked from 'package:tidyr':
##
##     expand, pack, unpack

## Registered S3 method overwritten by 'mosaic':
##   method                                from
##   fortify.SpatialPolygonsDataFrame ggplot2

##
## The 'mosaic' package masks several functions from core packages in order to add
## additional features. The original behavior of these functions should not be affected by
##
## Note: If you use the Matrix package, be sure to load it BEFORE loading mosaic.

##
## Attaching package: 'mosaic'

## The following object is masked from 'package:Matrix':
##
##     mean

## The following objects are masked from 'package:dplyr':
##
##     count, do, tally

## The following object is masked from 'package:purrr':
##
##     cross

## The following object is masked from 'package:ggplot2':
##
##     stat

## The following objects are masked from 'package:stats':
##
##     binom.test, cor, cor.test, cov, fivenum, IQR, median, prop.test,
##     quantile, sd, t.test, var

## The following objects are masked from 'package:base':
##
##     max, mean, min, prod, range, sample, sum

```

```
library(broom)
library(modelr)
```

```
##
## Attaching package: 'modelr'

## The following object is masked from 'package:broom':
##
##   bootstrap

## The following object is masked from 'package:mosaic':
##
##   resample

## The following object is masked from 'package:ggformula':
##
##   na.warn
```

```
library(car)
```

```
## Loading required package: carData

##
## Attaching package: 'car'

## The following objects are masked from 'package:mosaic':
##
##   deltaMethod, logit

## The following object is masked from 'package:dplyr':
##
##   recode

## The following object is masked from 'package:purrr':
##
##   some
```

```
library(knitr)
library(GGally)
```

```
## Registered S3 method overwritten by 'GGally':
##   method from
##   +.gg      ggplot2
```

```
##
## Attaching package: 'GGally'

## The following object is masked from 'package:dplyr':
##
##      nasa
```

```
library(dummies)
```

```
## dummies-1.5.6 provided by Decision Patterns
```

```
library(MASS)
```

```
##
## Attaching package: 'MASS'

## The following object is masked from 'package:dplyr':
##
##      select
```

```
library(foreign)
library(ISLR)
library(rsample)
library(rcfss)
```

```
##
## Attaching package: 'rcfss'

## The following object is masked from 'package:modelr':
##
##      mse
```

```
library(yardstick)
```

```
## For binary classification, the first factor level is assumed to be the event.
## Set the global option `yardstick.event_first` to `FALSE` to change this.
```

```
##
## Attaching package: 'yardstick'

## The following objects are masked from 'package:modelr':
##
##      mae, mape, rmse

## The following object is masked from 'package:readr':
##
##      spec
```

```
library(tree)
```

```
## Registered S3 method overwritten by 'tree':  
##   method      from  
##   print.tree cli
```

```
library(e1071)  
options(width=70, digits=4, scipen=8)  
knitr::opts_chunk$set(size='small') # Set the default R output size a bit smaller
```

Problem Set 3: Trees & Machines

Fork the repository

Fork the repository for the problem set 3, `problem-set-3` (<https://github.com/macss-m120/problem-set-3>). *Remember, all final submissions should be a single rendered PDF with code produced in-line.* Also, don't forget to **open the pull request** once you've committed your final submission to your forked repository. It needs to be merged back into the course master branch to be considered "submitted". See the syllabus for details.

The Data

ANES 2008

For the first part of this problem set, you will again be using the ANES data to predict feelings toward Joe Biden, but this time by growing and tuning several decision trees.

As a refresher, the `nes2008.csv` data contains a paired down selection of features from the full 2008 American National Election Studies survey. These data will allow you to test competing factors that may influence attitudes towards Joe Biden.

OJ Data

For the second part of the problem set on support vector machines, you will use the OJ data from the ISLR package, which contains 1070 purchases where the customer data related to orange juice purchases. A number of characteristics of the customer and product are recorded. Your goal here is to find a classification solution by tuning support vector classifiers that optimally predict whether customers buy either Citrus Hill (CH) or Minute Maid (MM) Orange Juice.

Important Note: Recall, there are many R and Python packages that allow you to complete the problem set in a variety of ways. I care less about your coding approach and selected packages, and significantly more about your results. Thus, feel free to use any programming language, and also any package/approach in your selected language.

Decision Trees

1. Set up the data and store some things for later use:

- Set seed
- Load the data
- Store the total number of features minus the biden feelings in object p
- Set λ (shrinkage/learning rate) range from 0.0001 to 0.04, by 0.001

```
set.seed(123)
nes2008 <- read.csv('nes2008.csv', header=T, sep=',')
p = 5
lamda=seq(from=0.0001,to=0.04,by=0.001)
```

2. (10 points) Create a training set consisting of 75% of the observations, and a test set with all remaining obs. **Note:** because you will be asked to loop over multiple λ values below, these training and test sets should only be integer values corresponding with row IDs in the data. This is a little tricky, but think about it carefully. If you try to set the training and testing sets as before, you will be unable to loop below.

```
set.seed(1234)
num =seq(from=1,to=1807,by=1)
train=sample(1:nrow(nes2008),1355)
test=rep("NA",times=452)
j=1
for(i in 1:length(num)){
  if(is.element(num[i],train)==FALSE){
    test[j]=num[i]
    j=j+1
  }
}
test=sample(test)
```

3. (15 points) Create empty objects to store training and testing MSE, and then write a loop to perform boosting on the training set with 1,000 trees for the pre-defined range of values of the shrinkage parameter, λ . Then, plot the training set and test set MSE across shrinkage values.

```
set.seed(1234)

library(gbm)
```

```
## Loaded gbm 2.1.5
```


```

length.lamda <- length(lamda)
train.errors <- rep(NA, length.lamda)
test.errors <- rep(NA, length.lamda)

for (i in 1:length.lamda) {
  boost.nes <- gbm(biden ~ .,
                    data=nes2008[train,],
                    distribution="gaussian",
                    n.trees=1000,
                    shrinkage=lamda[i],
                    interaction.depth = 4)
  train.pred <- predict(boost.nes, newdata=nes2008[train,],n.trees=1000)
  test.pred <- predict(boost.nes, newdata=nes2008[test,], n.trees=1000)
  train.errors[i] <- mean((nes2008[train,]$biden - train.pred)^2)
  test.errors[i] <- mean((nes2008[test,]$biden - test.pred)^2)
}

plot(lamda, train.errors, type="b",
     xlab="Shrinkage", ylab="Train MSE",
     col="blue", pch=20, bty = "n")


```



```

plot(lamda, test.errors, type="b",
     xlab="Shrinkage", ylab="Test MSE",
     col="blue", pch=20, bty = "n")

```



4. (10 points) The test MSE values are insensitive to some precise value of λ as long as its small enough. Update the boosting procedure by setting λ equal to 0.01 (but still over 1000 trees). Report the test MSE and discuss the results. How do they compare?

```
set.seed(200)
boost.nes <- gbm(biden ~ .,
                 data=nes2008[train,],
                 distribution="gaussian",
                 n.trees=1000,
                 shrinkage=0.01,
                 interaction.depth = 4)
test.pred <- predict(boost.nes, newdata=nes2008[test,], n.trees=1000)
mean((nes2008[test,]$biden - test.pred)^2)
```

```
## [1] 353.5
```

5. (10 points) Now apply bagging to the training set. What is the test set MSE for this approach?

```
set.seed(300)
library(randomForest)
```

```
## randomForest 4.6-14
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##
```

```
## Attaching package: 'randomForest'
```



```
## The following object is masked from 'package:dplyr':
```

```
##
```

```
##      combine
```

```
## The following object is masked from 'package:ggplot2':
```

```
##
```

```
##      margin
```

```
## The following object is masked from 'package:gridExtra':
```

```
##
```

```
##      combine
```

```
bag.nes <- randomForest(biden ~ ., data=nes2008[train,], ntree=1000, importance=TRUE)
bag.pred <- predict(bag.nes, nes2008[test,])
mean((nes2008[test,]$biden - bag.pred)^2)
```

```
## [1] 356.4
```

6. (10 points) Now apply random forest to the training set. What is the test set MSE for this approach?

```
set.seed(400)
rf.nes <- randomForest(biden ~ ., data=nes2008[train,], ntree=1000, importance=TRUE)
rf.pred <- predict(rf.nes, nes2008[test,])
mean((nes2008[test,]$biden - rf.pred)^2)
```

```
## [1] 357.6
```

7. (5 points) Now apply linear regression to the training set. What is the test set MSE for this approach?

```
set.seed(500)
biden_fu <- lm(biden ~ female + age + educ + dem + rep, data = nes2008[train,])
(test_mse <- augment(biden_fu, newdata = nes2008[test,]) %>%
  mse(truth = biden, estimate = .fitted))
```

```
## # A tibble: 1 x 3
```

```
##   .metric .estimator .estimate
```

```
##   <chr>    <chr>         <dbl>
```

```
## 1 mse      standard      345.
```

8. (5 points) Compare test errors across all fits. Discuss which approach generally fits best and how you concluded this.

#The boost(lamda = 0.01) MSE is 353.5. The bagging MSE is 356.4. The random forest MSE is 357.6. The linear regression MSE is 345.2. In this case, our linear regression fits the best since the MSE has the smallest value.

Support Vector Machines

1. Create a training set with a random sample of size 800, and a test set containing the remaining observations.

```
set.seed(3)
tr <- sample(1:nrow(OJ), 800)
oj.train <- OJ[tr,]
oj.test <- OJ[-tr,]
```

2. (10 points) Fit a support vector classifier to the training data with `cost = 0.01`, with `Purchase` as the response and *all* other features as predictors. Discuss the results.

```
svm_linear <- svm(Purchase ~ ., data = oj.train,
                  kernel = 'linear',
                  cost = 0.01)
summary(svm_linear)
```

```
##
## Call:
## svm(formula = Purchase ~ ., data = oj.train, kernel = "linear",
##      cost = 0.01)
##
##
## Parameters:
##   SVM-Type:  C-classification
## SVM-Kernel:  linear
##      cost:   0.01
##
## Number of Support Vectors:  434
##
## ( 216 218 )
##
##
## Number of Classes:  2
##
## Levels:
##  CH MM
```

3. (5 points) Display the confusion matrix for the classification solution, and also report both the training and test set error rates.

```
require(caret)
```

```
## Loading required package: caret
```

```
##
## Attaching package: 'caret'

## The following objects are masked from 'package:yardstick':
##
##   precision, recall

## The following object is masked from 'package:mosaic':
##
##   dotPlot

## The following object is masked from 'package:purrr':
##
##   lift
```

```
oj.pred <- predict(svm_linear, oj.test)
table(predicted = oj.pred, true = oj.test$Purchase)
```

```
##           true
## predicted  CH  MM
##          CH 139  22
##          MM  24  85
```

```
postResample(predict(svm_linear, oj.train), oj.train$Purchase)
```

```
## Accuracy    Kappa
##   0.8263     0.6294
```

```
postResample(predict(svm_linear, oj.test), oj.test$Purchase)
```

```
## Accuracy    Kappa
##   0.8296     0.6451
```

4. (10 points) Find an optimal cost in the range of 0.01 to 1000 (specific range values can vary; there is no set vector of range values you must use).

```
svm_linear_tune <-
  train(Purchase ~ ., data = oj.train, method = 'svmLinear2', trControl = trainControl(metho
svm_linear_tune
```

```
## Support Vector Machines with Linear Kernel
##
## 800 samples
## 17 predictor
## 2 classes: 'CH', 'MM'
##
## Pre-processing: centered (17), scaled (17)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 720, 720, 720, 720, 720, 720, ...
## Resampling results across tuning parameters:
##
## cost      Accuracy  Kappa
## 0.010     0.8237     0.6229
## 5.273     0.8275     0.6323
## 10.535    0.8225     0.6218
## 15.798    0.8250     0.6277
## 21.061    0.8225     0.6228
## 26.323    0.8225     0.6232
## 31.586    0.8225     0.6232
## 36.848    0.8225     0.6230
## 42.111    0.8225     0.6230
## 47.374    0.8225     0.6230
## 52.636    0.8225     0.6230
## 57.899    0.8225     0.6225
## 63.162    0.8225     0.6225
## 68.424    0.8237     0.6244
## 73.687    0.8237     0.6244
## 78.949    0.8237     0.6244
## 84.212    0.8237     0.6244
## 89.475    0.8237     0.6244
## 94.737    0.8237     0.6244
## 100.000   0.8237     0.6244
##
## Accuracy was used to select the optimal model using the
## largest value.
## The final value used for the model was cost = 5.273.
```

5. (10 points) Compute the optimal training and test error rates using this new value for cost. Display the confusion matrix for the classification solution, and also report both the training and test set error rates. How do the error rates compare? Discuss the results in substantive terms (e.g., how well did your optimally tuned classifier perform? etc.)

```
postResample(predict(svm_linear_tune, oj.train), oj.train$Purchase)
```

```
## Accuracy    Kappa
## 0.8387      0.6577
```

```
postResample(predict(svm_linear_tune, oj.test), oj.test$Purchase)
```

```
## Accuracy    Kappa  
##    0.8370    0.6627
```

```
oj.pred2 <- predict(svm_linear_tune, oj.test)  
table(predicted = oj.pred, true = oj.test$Purchase)
```

```
##           true  
## predicted  CH  MM  
##           CH 139 22  
##           MM  24 85
```

#The accuracy of the optimal model gets better. The optimally tuned classifier performs pretty good with the type I error rate equals to 22 and type II error rate equals to 24.